

# Data Manipulation with R

February 24, 2018

# Introduction to R Data Processing Package “dplyr”

Package “**dplyr**” in R focussed on tools for working with data frames

- ▶ dplyr provides abstractions for basic data manipulation operations (called verbs)
- ▶ **Verbs** can be combined to achieve complicated data manipulation results using a series of simple data processing steps (by building a data manipulation pipeline)
- ▶ The approach is familiar to those who use UNIX/Linux and the “**dotadiw**” philosophy: Do One Thing and Do It Well

# dplyr's Verbs for Data Processing

The verbs are:

- ▶ filter
- ▶ arrange
- ▶ select
- ▶ distinct
- ▶ mutate
- ▶ summarise

Data sets and R Code is available

<https://github.com/kiat/R-Examples>

```
#install.pacakges('dplyr')  
library(dplyr)  
delay.dat.houston <- read.csv("./Datasets/HoustonAirline.csv",  
                             header=TRUE,  
                             stringsAsFactors = FALSE)  
# tbl_df allows for nice printing  
delay.dat.houston <- tbl_df(delay.dat.houston)
```

# Take a look

```
> delay.dat.houston
# A tibble: 241,105 x 29
   Year Month DayofMonth DayOfWeek DepTime CRSDepTime ArrTime CRSArrTime
  <int> <int>    <int>    <int>    <int>    <int>    <int>    <int>
1  2008     1         4         5    1910      1910    2025      2025
2  2008     1         4         5    1345      1345    1453      1500
3  2008     1         4         5     736       735     839       850
4  2008     1         4         5    1603      1550    1647      1635
5  2008     1         4         5    2105      2105    2151      2150
6  2008     1         4         5     635       635     716       720
7  2008     1         4         5    1331      1330    1411      1415
8  2008     1         4         5    1850      1850    1936      1935
9  2008     1         4         5     956       1000    1038      1045
10 2008     1         4         5     823       805     906       850
# ... with 241,095 more rows, and 21 more variables: UniqueCarrier <chr>,
#   FlightNum <int>, TailNum <chr>, ActualElapsedTime <int>,
#   CRSElapsedTime <int>, AirTime <int>, ArrDelay <int>, DepDelay <int>,
#   Origin <chr>, Dest <chr>, Distance <int>, TaxiIn <int>, TaxiOut <int>,
#   Cancelled <int>, CancellationCode <chr>, Diverted <int>, CarrierDelay <int>,
#   WeatherDelay <int>, NASDelay <int>, SecurityDelay <int>,
#   LateAircraftDelay <int>
```

# Take a look

```
> airport.dat
# A tibble: 3,376 x 7
```

|    | iata<br><chr> | airport<br><chr>     | city<br><chr>    | state<br><chr> | country<br><chr> | lat<br><dbl> | long<br><dbl> |
|----|---------------|----------------------|------------------|----------------|------------------|--------------|---------------|
| 1  | 00M           | Thigpen              | Bay Springs      | MS             | USA              | 31.95376     | -89.23450     |
| 2  | 00R           | Livingston Municipal | Livingston       | TX             | USA              | 30.68586     | -95.01793     |
| 3  | 00V           | Meadow Lake          | Colorado Springs | CO             | USA              | 38.94575     | -104.56989    |
| 4  | 01G           | Perry-Warsaw         | Perry            | NY             | USA              | 42.74135     | -78.05208     |
| 5  | 01J           | Hilliard Airpark     | Hilliard         | FL             | USA              | 30.68801     | -81.90594     |
| 6  | 01M           | Tishomingo County    | Belmont          | MS             | USA              | 34.49167     | -88.20111     |
| 7  | 02A           | Gragg-Wade           | Clanton          | AL             | USA              | 32.85049     | -86.61145     |
| 8  | 02C           | Capitol              | Brookfield       | WI             | USA              | 43.08751     | -88.17787     |
| 9  | 02G           | Columbiana County    | East Liverpool   | OH             | USA              | 40.67331     | -80.64141     |
| 10 | 03D           | Memphis Memorial     | Memphis          | MO             | USA              | 40.44726     | -92.22696     |

```
# ... with 3,366 more rows
```

# Variable Description

## Variable descriptions

| Name                 | Description   |
|----------------------|---|
| 1 Year               | 1987-2008   |
| 2 Month              | 1-12  |
| 3 DayofMonth         | 1-31  |
| 4 DayOfWeek          | 1 (Monday) - 7 (Sunday)   |
| 5 DepTime            | actual departure time (local, hhmm)                                       |
| 6 CRSDEPTime         | scheduled departure time (local, hhmm)                                    |
| 7 ArrTime            | actual arrival time (local, hhmm)   |
| 8 CRSARRTime         | scheduled arrival time (local, hhmm)                                      |
| 9 UniqueCarrier      | <a href="#"><u>unique carrier code</u></a>                                |
| 10 FlightNum         | flight number   |
| 11 TailNum           | plane tail number   |
| 12 ActualElapsedTime | in minutes  |
| 13 CRSElapsedTime    | in minutes  |
| 14 AirTime           | in minutes  |
| 15 ArrDelay          | arrival delay, in minutes   |
| 16 DepDelay          | departure delay, in minutes   |
| 17 Origin            | origin <a href="#"><u>ATA airport code</u></a>                            |
| 18 Dest              | destination <a href="#"><u>ATA airport code</u></a>                       |
| 19 Distance          | in miles  |
| 20 TaxiIn            | taxi in time, in minutes  |
| 21 TaxiOut           | taxi out time in minutes  |
| 22 Cancelled         | was the flight cancelled?   |
| 23 CancellationCode  | reason for cancellation (A = carrier, B = weather, C = NAS, D = security) |
| 24 Diverted          | 1 = yes, 0 = no   |
| 25 CarrierDelay      | in minutes  |
| 26 WeatherDelay      | in minutes  |
| 27 NASDelay          | in minutes  |
| 28 SecurityDelay     | in minutes  |
| 29 LateAircraftDelay | in minutes  |

# Filter - Filtering the Data

- ▶ filter is probably the most familiar verb
- ▶ filter is dplyr's version of R's subset() function
- ▶ filter returns all rows (observations) for which a logical condition holds



# Filter - Inputs and Outputs

- ▶ Inputs: `data.frame` and logical expressions
- ▶ Output: `data.frame`
- ▶ All dplyr verbs behave similarly
- ▶ A `data.frame` is inputted, and a `data.frame` is outputted

# Example 1

```
# Find all flight which occurred in January
> filter(delay.dat.houston, Month==1)
# A tibble: 20,349 x 29
   Year Month DayOfMonth DayOfWeek DepTime CRSDepTime ArrTime CRSArrTime
  <int> <int>      <int>      <int>   <int>      <int>      <int>      <int>
1  2008     1         4         5    1910        1910        2025        2025
2  2008     1         4         5    1345        1345        1453        1500
3  2008     1         4         5     736         735         839         850
4  2008     1         4         5    1603        1550        1647        1635
5  2008     1         4         5    2105        2105        2151        2150
6  2008     1         4         5     635         635         716         720
7  2008     1         4         5    1331        1330        1411        1415
8  2008     1         4         5    1850        1850        1936        1935
9  2008     1         4         5     956        1000        1038        1045
10 2008     1         4         5     823         805         906         850
# ... with 20,339 more rows, and 21 more variables: UniqueCarrier <chr>,
# FlightNum <int>, TailNum <chr>, ActualElapsedTime <int>, CRSElapsedTime <int>,
# AirTime <int>, ArrDelay <int>, DepDelay <int>, Origin <chr>, Dest <chr>,
# Distance <int>, TaxiIn <int>, TaxiOut <int>, Cancelled <int>,
# CancellationCode <chr>, Diverted <int>, CarrierDelay <int>, WeatherDelay <int>,
# NASDelay <int>, SecurityDelay <int>, LateAircraftDelay <int>
```

## Example 2

```
# Using airport data, find a list of iata abbreviations for houston texas airports
> filter(airport.dat, state=='TX', city=='Houston')
# A tibble: 8 x 7
```

|   | iata<br><chr> | airport<br><chr>             | city<br><chr> | state<br><chr> | country<br><chr> | lat<br><dbl> | long<br><dbl> |
|---|---------------|------------------------------|---------------|----------------|------------------|--------------|---------------|
| 1 | DWH           | David Wayne Hooks Memorial   | Houston       | TX             | USA              | 30.06186     | -95.55278     |
| 2 | EFD           | Ellington                    | Houston       | TX             | USA              | 29.60733     | -95.15875     |
| 3 | HOU           | William P Hobby              | Houston       | TX             | USA              | 29.64542     | -95.27889     |
| 4 | IAH           | George Bush Intercontinental | Houston       | TX             | USA              | 29.98047     | -95.33972     |
| 5 | IWS           | West Houston                 | Houston       | TX             | USA              | 29.81819     | -95.67261     |
| 6 | LVJ           | Clover                       | Houston       | TX             | USA              | 29.52131     | -95.24217     |
| 7 | SGR           | Sugar Land Municipal/Hull    | Houston       | TX             | USA              | 29.62225     | -95.65653     |
| 8 | SPX           | Houston-Gulf                 | Houston       | TX             | USA              | 29.50836     | -95.05133     |

# Introduction to R Data Processing Package “dplyr”

Package “**dplyr**” in R focussed on tools for working with data frames

- ▶ Find the subset of flight departing from Hobby Airport “**HOU**” for which the Actual Elapsed Time was greater than the CRS Elapsed Time ( $\text{ActualElapsedTime} > \text{CRSElapsedTime}$ )
- ▶ Find the subset of flights departing on the weekend.

# R Command

```
# Find the subset of flight departing from  
# Hobby Airport "HOU" for which the Actual  
# Elapsed Time was greater than the CRS Elapsed Time.  
> filter(delay.dat.houston,  
         Origin == 'HOU', # iata code for Hobby  
         ActualElapsedTime > CRSElapsedTime)
```

# R Command

```
# Find the subset of flights departing on the weekend.  
> filter(delay.dat.houston, DayOfWeek == 6 | DayOfWeek == 7)  
  
# another alternative  
> filter(delay.dat.houston, DayOfWeek %in% c(6,7))
```

# arrange

- ▶ `arrange`, like `filter`, operates on `data.frame` rows
- ▶ `arrange` is used for sorting `data.frame` rows w.r.t. a given column(s)

# arrange

```
> arrange(delay.dat.houston, DayofMonth)
# A tibble: 241,105 x 29
   Year Month DayofMonth DayOfWeek DepTime CRSDepTime ArrTime CRSArrTime
  <int> <int>    <int>    <int>    <int>    <int>    <int>    <int>
1  2008     1         1         2    1531      1525    1626    1622
2  2008     1         1         2    1848      1850    2022    2025
3  2008     1         1         2    1024      1025    1353    1352
4  2008     1         1         2     707       705     818     822
5  2008     1         1         2    1047      1045    1423    1415
6  2008     1         1         2    1110      1110    1237    1240
7  2008     1         1         2    1653      1655    2038    2058
8  2008     1         1         2    2013      1950    2335    2319
9  2008     1         1         2    1212      1220    1454    1512
10 2008     1         1         2    1021      1020    1136    1132
# ... with 241,095 more rows, and 21 more variables: UniqueCarrier <chr>,
#   FlightNum <int>, TailNum <chr>, ActualElapsedTime <int>, CRSElapsedTime <int>,
#   AirTime <int>, ArrDelay <int>, DepDelay <int>, Origin <chr>, Dest <chr>,
#   Distance <int>, TaxiIn <int>, TaxiOut <int>, Cancelled <int>,
#   CancellationCode <chr>, Diverted <int>, CarrierDelay <int>, WeatherDelay <int>,
#   NASDelay <int>, SecurityDelay <int>, LateAircraftDelay <int>
```



# arrange

```
> arrange(delay.dat.houston, desc(Month), desc(DayofMonth))
# A tibble: 241,105 x 29
   Year Month DayofMonth DayOfWeek DepTime CRSDepTime ArrTime CRSArrTime
  <int> <int>    <int>    <int>    <int>    <int>    <int>    <int>
1  2008    12         31         3      707         705      810      815
2  2008    12         31         3     1256        1245     1355     1400
3  2008    12         31         3     1553        1550     1632     1635
4  2008    12         31         3     1801        1750     1841     1835
5  2008    12         31         3     1101        1055     1141     1140
6  2008    12         31         3     1325        1315     1408     1400
7  2008    12         31         3      948         950     1113     1125
8  2008    12         31         3     1555        1555     1719     1730
9  2008    12         31         3     1952        1955     2124     2135
10 2008    12         31         3     1755        1720     1936     1910
# ... with 241,095 more rows, and 21 more variables: UniqueCarrier <chr>,
#   FlightNum <int>, TailNum <chr>, ActualElapsedTime <int>, CRSElapsedTime <int>,
#   AirTime <int>, ArrDelay <int>, DepDelay <int>, Origin <chr>, Dest <chr>,
#   Distance <int>, TaxiIn <int>, TaxiOut <int>, Cancelled <int>,
#   CancellationCode <chr>, Diverted <int>, CarrierDelay <int>, WeatherDelay <int>,
#   NASDelay <int>, SecurityDelay <int>, LateAircraftDelay <int>
```

# arrange

```
> arrange(delay.dat.houston, desc(Month), desc(DayofMonth))
# A tibble: 241,105 x 29
   Year Month DayofMonth DayOfWeek DepTime CRSDepTime ArrTime CRSArrTime
  <int> <int>    <int>    <int>    <int>    <int>    <int>    <int>
1  2008    12         31         3      707         705      810      815
2  2008    12         31         3     1256        1245     1355     1400
3  2008    12         31         3     1553        1550     1632     1635
4  2008    12         31         3     1801        1750     1841     1835
5  2008    12         31         3     1101        1055     1141     1140
6  2008    12         31         3     1325        1315     1408     1400
7  2008    12         31         3      948         950     1113     1125
8  2008    12         31         3     1555        1555     1719     1730
9  2008    12         31         3     1952        1955     2124     2135
10 2008    12         31         3     1755        1720     1936     1910
# ... with 241,095 more rows, and 21 more variables: UniqueCarrier <chr>,
# FlightNum <int>, TailNum <chr>, ActualElapsedTime <int>, CRSElapsedTime <int>,
# AirTime <int>, ArrDelay <int>, DepDelay <int>, Origin <chr>, Dest <chr>,
# Distance <int>, TaxiIn <int>, TaxiOut <int>, Cancelled <int>,
# CancellationCode <chr>, Diverted <int>, CarrierDelay <int>, WeatherDelay <int>,
# NASDelay <int>, SecurityDelay <int>, LateAircraftDelay <int>
```

# select

- ▶ select is like filter but for columns
- ▶ select is used for keeping/dropping a subset of variables/columns

# R Command

Try out the following examples using select

```
select(delay.dat.houston, Year, Month, DayofMonth)
```

```
select(delay.dat.houston, Year:DayofMonth)
```

```
select(delay.dat.houston, -(Year:DayofMonth))
```

# select

Here we use the contains helper:

```
> select(delay.dat.houston, contains('Dep'))
```

```
# A tibble: 241,105 x 3
```

|    | DepTime | CRSDepTime | DepDelay |
|----|---------|------------|----------|
|    | <int>   | <int>      | <int>    |
| 1  | 1910    | 1910       | 0        |
| 2  | 1345    | 1345       | 0        |
| 3  | 736     | 735        | 1        |
| 4  | 1603    | 1550       | 13       |
| 5  | 2105    | 2105       | 0        |
| 6  | 635     | 635        | 0        |
| 7  | 1331    | 1330       | 1        |
| 8  | 1850    | 1850       | 0        |
| 9  | 956     | 1000       | -4       |
| 10 | 823     | 805        | 18       |

```
# ... with 241,095 more rows
```

# select helper

Create a select statement using

- ▶ `one_of` helper
- ▶ `ends_with` helper

## select helper

```
select(delay.dat.houston,  
        one_of('UniqueCarrier',  
               'FlightNum'))
```

```
select(delay.dat.houston,  
        ends_with('Time'))
```

# distinct

- ▶ distinct finds unique values of a variable
- ▶ distinct returns the first observation/row containing each value



# distinct

```
> distinct(delay.dat.houston, Month)
```

```
# A tibble: 12 x 1
```

```
  Month
```

```
  <int>
```

|    |    |
|----|----|
| 1  | 1  |
| 2  | 2  |
| 3  | 3  |
| 4  | 4  |
| 5  | 5  |
| 6  | 6  |
| 7  | 7  |
| 8  | 8  |
| 9  | 9  |
| 10 | 10 |
| 11 | 11 |
| 12 | 12 |

# distinct

```
> distinct(delay.dat.houston, Month, .keep_all=TRUE)
# A tibble: 12 x 29
   Year Month DayOfMonth DayOfWeek DepTime CRSDepTime ArrTime CRSArrTime
  <int> <int>      <int>      <int>   <int>      <int>      <int>      <int>
1  2008     1         4         5    1910        1910        2025        2025
2  2008     2         3         7     758         800         903         915
3  2008     3         3         1     800         800         920         915
4  2008     4         4         5     900         900        1027        1010
5  2008     5         4         7     857         900        1008        1010
6  2008     6         3         2    1951        1935        2050        2040
7  2008     7         3         4    1935        1935        2032        2040
8  2008     8         3         7    1940        1935        2049        2040
9  2008     9         3         3     804         805         857         910
10 2008    10         3         5     715         720         828         845
11 2008    11         4         2    1834        1825        1933        1935
12 2008    12         3         3    1845        1825        1958        1935
# ... with 21 more variables: UniqueCarrier <chr>, FlightNum <int>, TailNum <chr>,
#   ActualElapsedTime <int>, CRSElapsedTime <int>, AirTime <int>, ArrDelay <int>,
#   DepDelay <int>, Origin <chr>, Dest <chr>, Distance <int>, TaxiIn <int>,
#   TaxiOut <int>, Cancelled <int>, CancellationCode <chr>, Diverted <int>,
#   CarrierDelay <int>, WeatherDelay <int>, NASDelay <int>, SecurityDelay <int>,
#   LateAircraftDelay <int>
```

# distinct

```
> distinct(delay.dat.houston, Month, DayOfWeek)
```

```
# A tibble: 84 x 2
```

```
  Month DayOfWeek
```

```
  <int>    <int>
```

```
1     1         5
```

```
2     1         6
```

```
3     1         7
```

```
4     1         1
```

```
5     1         2
```

```
6     1         3
```

```
7     1         4
```

```
8     2         7
```

```
9     2         1
```

```
10    2         2
```

```
# ... with 74 more rows
```

## Combination of verbs

You can combine distinct with the select verb from previous.  
What do you think the following will do?

```
select(  
  distinct(  
    arrange(  
      filter(delay.dat.houston, DayOfWeek==6),  
      desc(ActualElapsedTime)),  
    UniqueCarrier, .keep_all = TRUE),  
    UniqueCarrier, ActualElapsedTime)
```

Reading from the inside out we can see it:

- ▶ Only considers flights departing on Saturday
- ▶ Arranges these by ActualElapsedTime in decrease order
- ▶ Selects the first row for each carrier
- ▶ In total this gives the largest ActualElapsedTime for Saturday departing flights for each carrier.
- ▶ distinct returns the first observation/row containing each value

# Combination of verbs

We can do the previous example with the chaining

```
delay.dat.houston %>%  
  filter(DayOfWeek == 6) %>%  
  arrange(desc(ActualElapsedTime)) %>%  
  distinct(UniqueCarrier, .keep_all=TRUE) %>%  
  select(UniqueCarrier, ActualElapsedTime)
```

Chain together the verbs we've seen so far to:

- ▶ Find a list of Origin Airports
- ▶ Find a list of (Origin, Destination) pairs
- ▶ Find the Origin airport which had the largest departure delay in the month of January
- ▶ Find the largest departure delay for each carrier for each month

# Combination of verbs

```
# Find a list of the distinct Origin airports
```

```
delay.dat.houston %>%  
  distinct(Origin)
```

```
# Find a list of distinct (Origin, Dest) pairs
```

```
delay.dat.houston %>%  
  distinct(Origin, Dest)
```

```
# Origin airport with largest January departure delay
```

```
delay.dat.houston %>%  
  filter(Month==1) %>%  
  arrange(desc(DepDelay)) %>%  
  select(Month, Origin, DepDelay) %>%  
  distinct(Origin, .keep_all = TRUE)
```



# Combination of verbs

```
# largest departure delay for each carrier for each month
delay.dat.houston %>%
  arrange(Month,desc(DepDelay)) %>%
  select(Month,UniqueCarrier,DepDelay) %>%
  distinct(Month,UniqueCarrier,.keep_all=TRUE)
```

Two verbs: mutate and summarise

- ▶ mutate allows us to create new variables

summarise:

- ▶ summarise let's us compute summary statistics on groups of data
- ▶ summarise is used in conjunction with the group by verb

# summarise

```
# Basic example with no grouping
> delay.dat.houston %>%
  summarise(MeanDistance = mean(Distance,na.rm=TRUE))
# A tibble: 1 x 1
  MeanDistance
      <dbl>
1      778.5913
```

# summarise

```
# With grouping
# n() is dplyr function counts # obs in each group
> delay.dat.houston %>%
+   group_by(UniqueCarrier) %>%
+   summarise(
+     MeanDistance=mean(Distance,na.rm=TRUE),
+     NFlights = n())
# A tibble: 17 x 3
```

|    | UniqueCarrier | MeanDistance | NFlights |
|----|---------------|--------------|----------|
|    | <chr>         | <dbl>        | <int>    |
| 1  | 9E            | 630.9294     | 2721     |
| 2  | AA            | 586.7512     | 4325     |
| 3  | B6            | 1428.0000    | 944      |
| 4  | CO            | 1055.0753    | 85642    |
| 5  | DL            | 690.3982     | 1517     |
| 6  | EV            | 704.0464     | 194      |
| 7  | F9            | 861.0000     | 846      |
| 8  | FL            | 696.0000     | 1792     |
| 9  | MQ            | 247.0000     | 2425     |
| 10 | NW            | 1013.0745    | 1598     |
| 11 | OH            | 912.1431     | 1013     |
| 12 | OO            | 1007.0786    | 2595     |
| 13 | UA            | 1019.8512    | 2325     |
| 14 | US            | 965.7900     | 1924     |
| 15 | WN            | 562.0526     | 48968    |
| 16 | XE            | 611.3961     | 80194    |
| 17 | YV            | 991.3463     | 2082     |

We could also redo our previous example, finding the largest departure delay for each carrier for each month

```
> delay.dat.houston %>%  
+   group_by(Month, UniqueCarrier) %>%  
+   summarise(MaxDepDelay = max(DepDelay, na.rm=TRUE)) %>%  
+   head(5)
```

# A tibble: 5 x 3

# Groups: Month [1]

|   | Month | UniqueCarrier | MaxDepDelay |
|---|-------|---------------|-------------|
|   | <int> | <chr>         | <dbl>       |
| 1 | 1     | 9E            | 356         |
| 2 | 1     | AA            | 234         |
| 3 | 1     | B6            | 183         |
| 4 | 1     | C0            | 475         |
| 5 | 1     | DL            | 131         |

- ▶ For each carrier plot the average Departure delay for each month.
- ▶ Do you notice anything strange? What might be the cause?
- ▶ Hint: Use summarise and faceting
- ▶ Hint: For each carrier also plot the number of flights per month.

```
library(ggplot2)
delay.dat.houston %>%
  group_by(Month,UniqueCarrier) %>%
  summarise(
    Dep = mean(DepDelay,na.rm=TRUE)
  ) -> tmp
```

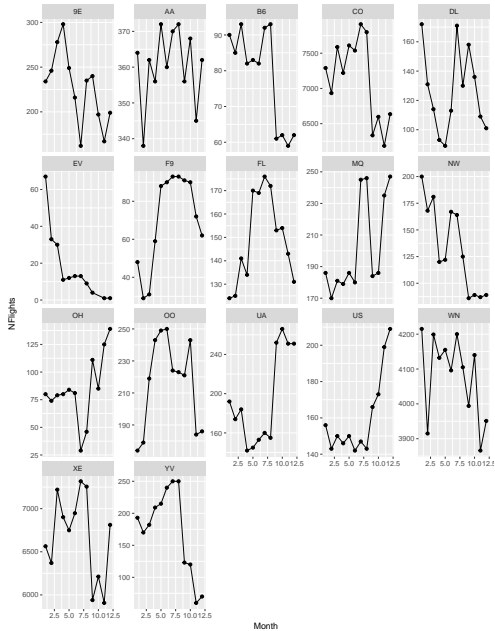
```
qplot(Month,Dep,data=tmp) +
  geom_line() +
  facet_wrap(~UniqueCarrier)
```





What could cause this? Try this:

```
delay.dat.houston %>%  
  group_by(Month,UniqueCarrier) %>%  
  summarise(  
    NFlights = n()  
  ) -> tmp  
  
qplot(Month,NFlights,data=tmp) +  
  geom_line() +  
  facet_wrap(~UniqueCarrier,scale='free_y')
```



- ▶ Find the percent of flights cancelled for each carrier.
- ▶ Use summarise to get total number of flights for each carrier (UniqueCarrier) and the total number of cancelled flights
- ▶ Create a new variable PercentCancelled based on the results above
- ▶ Return a data.frame with only UniqueCarrier and PercentCancelled

```
delay.dat.houston %>%  
  group_by(UniqueCarrier) %>%  
  summarise(  
    NFlights = n(),  
    NCancelled = sum(Cancelled)) %>%  
  mutate(  
    PercentCancelled = (NCancelled/NFlights)*100) %>%  
  select(UniqueCarrier,  
    PercentCancelled)
```


# A tibble: 17 x 2

|    | UniqueCarrier | PercentCancelled |
|----|---------------|------------------|
|    | <chr>         | <dbl>            |
| 1  | 9E            | 3.601617         |
| 2  | AA            | 4.138728         |
| 3  | B6            | 3.283898         |
| 4  | CO            | 1.122113         |
| 5  | DL            | 2.834542         |
| 6  | EV            | 3.092784         |
| 7  | F9            | 1.418440         |
| 8  | FL            | 1.450893         |
| 9  | MQ            | 3.835052         |
| 10 | NW            | 1.251564         |
| 11 | OH            | 3.849951         |
| 12 | OO            | 2.581888         |
| 13 | UA            | 2.408602         |
| 14 | US            | 1.663202         |
| 15 | WN            | 2.783450         |
| 16 | XE            | 2.188443         |
| 17 | YV            | 3.073967         |

- ▶ For each Destination find the average Arrival and Departure delay; create associated variables AvgArrDel, AvgDepDel
- ▶ Plot AvgArrDel vs AvgDepDel for the three largest carriers (largest in terms of number of flights)
- ▶ Plot AvgArrDel vs AvgDepDel for all carriers. Use point size to indicate carrier size

```
delay.dat.houston %>%  
  group_by(UniqueCarrier) %>%  
  summarise(  
    Dep = mean(DepDelay,na.rm=TRUE),  
    Arr = mean(ArrDelay,na.rm=TRUE),  
    NFlights = n()  
  ) %>%  
  select(Dep,Arr,NFlights) -> tmp
```

```
qplot(Dep,  
      Arr,  
      data=tmp,  
      size=log(NFlights))+  
geom_abline(intercept=0,slope=1,color='red')
```





For our final dplyr stop we'll look at it's merging capabilities.  
Let's start by reading in some more toy datasets

```
people.info <- read.table('./Datasets/mergedata/PeopleInfo.csv',  
                           sep=',',  
                           header=TRUE)  
  
occup.info <-  
↪ read.table('./Datasets/mergedata/OccupationInfo.csv',  
              sep=',',  
              header=TRUE)
```

# People Dataset

```
> people.info
```

|   | ID   | Last    | DOB |
|---|------|---------|-----|
| 1 | 1718 | Jones   | 85  |
| 2 | 1817 | Smith   | 72  |
| 3 | 1558 | Wallace | 50  |
| 4 | 1742 | Marks   | 90  |

```
> occup.info
```

|   | ID   | Title      | Office |
|---|------|------------|--------|
| 1 | 1558 | Supervisor | 101    |
| 2 | 1718 | Clerk      | 110    |
| 3 | 2234 | Accountant | 502    |
| 4 | 943  | Doctor     | 409    |
| 5 | 1119 | Manager    | 404    |

# Basic Join

dplyr's basic merging functions are:

- ▶ **inner\_join** : return all rows from x where there are matching values in y, and all columns from x and y. If there are multiple matches between x and y, all combination of the matches are returned.
- ▶ **left\_join** : return all rows from x, and all columns from x and y. Rows in x with no match in y will have NA values in the new columns. If there are multiple matches between x and y, all combinations of the matches are returned.
- ▶ **right\_join** :

# INNER Join

Venn Diagram for Join Operation.




Figure: TableA INNER JOIN TableB

# Full LEFT OUTER Join

Full Left Outer Join.




Figure: TableA LEFT OUTER JOIN TableB

# Join

```
# What do you think the following snippets will do
# Try to guess before running, then run to confirm
left_join(people.info, occup.info)
right_join(people.info, occup.info)
inner_join(people.info, occup.info)

# Do the following return the same data set?
left_join(people.info, occup.info)
right_join(occup.info, people.info)

# Do you think this will work?
people.info %>% left_join(occup.info)
```

## Other Join

- ▶ **semi\_join** returns only lhs columns, and only for ids common to both
- ▶ **anti\_join** returns only lhs columns, and only for ids *\*not\** common to both
- ▶ **full\_join** returns all columns, for all ids, merging with inner/left/right when applicable

# LEFT OUTER JOIN Join

TableA LEFT OUTER JOIN TableB




Figure: TableA LEFT OUTER JOIN TableB



# FULL OUTER Join

TableA FULL OUTER JOIN TableB




Figure: TableA FULL OUTER JOIN TableB

## full join

TableA FULL OUTER JOIN TableB




Figure: TableA FULL OUTER JOIN TableB

# Join

```
semi_join(people.info, occup.info)  
anti_join(people.info, occup.info)  
full_join(people.info, occup.info)
```

# Join

Merge the airport and delay data so that we have state/city information regarding the destination

**Hint:** use `left_join` with `by=c("Dest" = "iata")`

# Join

```
delay.dat.houston %>%  
  left_join(airport.dat,  
            by=c("Dest" = "iata"))
```

# Join

Calculate the number of flights to each destination state

For each carrier, for which state do they have the largest average delay?

# Join

```
# one option
delay.dat.houston %>%
  left_join(airport.dat,
            by=c("Dest" = 'iata')) %>%
  group_by(state) %>%
  summarise(
    NFlights = n()
  ) %>%
  select(state, NFlights)
```

```
# A tibble: 41 x 2
```

```
  state NFlights
```

```
  <chr>   <int>
```

```
1    AK     206
```

```
2    AL   5778
```

```
3    AR   2911
```

```
4    AZ   7568
```

```
5    CA  17448
```

```
6    CO   7869
```

```
7    CT    120
```

```
8    FL  18951
```

```
9    GA   9533
```

```
10   HI    702
```

```
# ... with 31 more rows
```



```
delay.dat.houston %>%  
  left_join(airport.dat,  
            by=c("Dest" = 'iata')) %>%  
  group_by(UniqueCarrier, state) %>%  
  summarise(  
    AvgDelay = mean(DepDelay, na.rm=TRUE)  
  ) %>%  
  select(state, UniqueCarrier, AvgDelay) %>%  
  arrange(UniqueCarrier, desc(AvgDelay)) %>%  
  distinct(UniqueCarrier, .keep_all = TRUE)
```